

Build Survival Model: Random Survival Forest

Mingcheng Hu

Table of contents

Load Data	2
Random Survival Forest (RSF)	4
Variable Selection	4
Cross Validation to Select the Best Number of Features	7
Model Fitting	9

```
library(tidyverse)
library(survival)
library(randomForestSRC)
library(caret)
library(survcomp)
library(parallel)
library(doParallel)
library(mcprogress) # wrap mclapply with progress bar.
library(kableExtra) # include knitr automatically

source("/work/users/y/u/youkias/BIOS-Material/BIOS992/utlis/csv_utils.r")
# * Don't use setwd() for Quarto documents!
# setwd("/work/users/y/u/youkias/BIOS-Material/BIOS992/data")

adjust_type <- ifelse(exists("params"), params$adjust_type, "partial") #
  ↳ options: "minimal", "partial", "full"
impute_type <- ifelse(exists("params"), params$impute_type, "unimputed") #
  ↳ options: "unimputed", "imputed"
include_statin <- ifelse(exists("params"), params$include_statin, "no") #
  ↳ options: "yes", "no"

n_folds <- 10
set.seed(1234)
```

```
# string of parameters
adjust_type_str <- switch(adjust_type,
  minimal = "minimal",
  partial = "partial",
  full = "full"
)
print(paste0("Model Adjustment Type: ", adjust_type_str))
```

```
[1] "Model Adjustment Type: partial"
```

```
impute_type_str <- switch(impute_type,
  unimputed = "unimputed",
  imputed = "imputed"
)
print(paste0("Data Imputation Type: ", impute_type_str))
```

```
[1] "Data Imputation Type: unimputed"
```

Load Data

```
if (include_statin == "yes") {
  data_train <-
  ↪ read.csv(paste0("/work/users/y/u/youkias/BIOS-Material/BIOS992/data/train_data_",
  ↪ impute_type_str, "_statin.csv"),
    header = TRUE
  )
} else {
  data_train <-
  ↪ read.csv(paste0("/work/users/y/u/youkias/BIOS-Material/BIOS992/data/train_data_",
  ↪ impute_type_str, ".csv"),
    header = TRUE
  )
}

data_train <- data_train[, -1] # the first column is the index generated by
  ↪ sklearn
(dim(data_train))
```

```
[1] 28127    100
```

```
data <- select_subset(data_train, type = adjust_type)
(dim(data))
```

```
[1] 28127    75
```

```
colnames(data)
```

```
[1] "event"           "time"
[3] "HRV_MeanNN"      "HRV_SDNN"
[5] "HRV_RMSSD"       "HRV_SDSD"
[7] "HRV_CVNN"        "HRV_CVSD"
[9] "HRV_MedianNN"    "HRV_MadNN"
[11] "HRV_MCVNN"       "HRV_IQRNN"
[13] "HRV_SDRMSSD"     "HRV_Prc20NN"
[15] "HRV_Prc80NN"     "HRV_pNN50"
[17] "HRV_pNN20"       "HRV_MinNN"
[19] "HRV_MaxNN"       "HRV_HTI"
[21] "HRV_TINN"        "HRV_LF"
[23] "HRV_HF"          "HRV_VHF"
[25] "HRV_TP"          "HRV_LFHF"
[27] "HRV_LFn"         "HRV_HFn"
[29] "HRV_LnHF"        "HRV_SD1"
[31] "HRV_SD2"         "HRV_SD1SD2"
[33] "HRV_S"           "HRV_CSI"
[35] "HRV_CVI"         "HRV_CSI_Modified"
[37] "HRV_PIP"         "HRV_IALS"
[39] "HRV_PSS"         "HRV_PAS"
[41] "HRV_GI"          "HRV_SI"
[43] "HRV_AI"          "HRV_PI"
[45] "HRV_C1d"         "HRV_C1a"
[47] "HRV_SD1d"        "HRV_SD1a"
[49] "HRV_C2d"         "HRV_C2a"
[51] "HRV_SD2d"        "HRV_SD2a"
[53] "HRV_Cd"          "HRV_Ca"
[55] "HRV_SDNNd"       "HRV_SDNNa"
[57] "HRV_ApEn"        "HRV_ShanEn"
[59] "HRV_FuzzyEn"     "HRV_MSEn"
[61] "HRV_CMSEn"       "HRV_RCMSEn"
```

```

[63] "HRV_CD"                "HRV_HFD"
[65] "HRV_KFD"               "HRV_LZC"
[67] "HRV_DFA_alpha1"        "HRV_MFDFA_alpha1_Width"
[69] "HRV_MFDFA_alpha1_Peak" "HRV_MFDFA_alpha1_Mean"
[71] "HRV_MFDFA_alpha1_Max"  "HRV_MFDFA_alpha1_Delta"
[73] "HRV_MFDFA_alpha1_Asymmetry" "HRV_MFDFA_alpha1_Fluctuation"
[75] "HRV_MFDFA_alpha1_Increment"

```

```
data <- tibble::as_tibble(data)
```

```

# * It is very hard to compare the HR as different predictors are on
  ↳ different magnitudes, so we need to normalize them.
time_col <- data$time
event_col <- data$event
data <- data %>%
  select(-c(time, event)) %>%
  mutate(across(where(is.numeric), scale)) %>%
  mutate(
    time = time_col,
    event = event_col
  )

```

Note now the interpretation of HR is different! For example, if $HR=1.16$ for the predictor in the univariate model fitted using scaled data, it means that each standard deviation increase is associated with 16% higher risk of event.

```
# For RSF model, we don't need to exclude the missing values
```

Random Survival Forest (RSF)

Variable Selection

The `method` argument can be set to `vh` instead for variable hunting, which should be used for problems where the number of variables is substantially larger than the sample size.

```

n_cores <- min(parallel::detectCores() - 1, 32)
cl <- makeCluster(n_cores)
registerDoParallel(cl)

```

```
rsf_var_select <- var.select.rfsrc(Surv(time, event) ~ .,
  data = data,
  method = "md",
  seed = 1234,
  ntree = 200,
  parallel = TRUE
) # minimal depth variable selection
```

running forests ...
minimal depth variable selection ...

```
-----
family           : surv
var. selection   : Minimal Depth
conservativeness : medium
x-weighting used? : TRUE
dimension        : 73
sample size      : 26782
ntree            : 200
nsplit          : 10
mtry            : 25
nodesize        : 2
refitted forest  : FALSE
model size       : 73
depth threshold  : 15.4073
PE (true OOB)    : 44.7123
```

Top variables:

	depth	vimp
HRV_Prc20NN	2.005	0.000
HRV_CD	3.375	0.004
HRV_MaxNN	3.480	0.016
HRV_KFD	3.760	0.004
HRV_SI	4.045	0.002
HRV_ApEn	4.420	0.002
HRV_MedianNN	4.600	-0.001
HRV_LZC	4.695	-0.004
HRV_PIP	4.705	0.000
HRV_IALS	4.810	0.000

HRV_HTI	4.850	-0.001
HRV_PI	4.890	0.001
HRV_pNN50	4.905	0.002
HRV_MadNN	4.915	0.003
HRV_MeanNN	4.950	0.008
HRV_ShanEn	5.020	-0.001
HRV_MinNN	5.115	0.007
HRV_TINN	5.165	0.004
HRV_PAS	5.175	0.003
HRV_SD2a	5.175	0.009
HRV_IQRNN	5.185	-0.003
HRV_GI	5.240	0.001
HRV_C2d	5.370	0.002
HRV_CMSEn	5.375	0.003
HRV_RCMSEn	5.415	-0.001
HRV_MFDFA_alpha1_Peak	5.420	0.002
HRV_MFDFA_alpha1_Asymmetry	5.460	-0.001
HRV_Prc80NN	5.475	0.002
HRV_MCVNN	5.505	0.000
HRV_pNN20	5.570	0.003
HRV_PSS	5.600	0.000
HRV_HFn	5.635	0.002
HRV_SDRMSSD	5.685	0.002
HRV_VHF	5.700	0.003
HRV_AI	5.750	0.003
HRV_LF	5.755	0.002
HRV_MFDFA_alpha1_Max	5.755	0.002
HRV_C1a	5.775	0.001
HRV_FuzzyEn	5.800	0.000
HRV_MFDFA_alpha1_Fluctuation	5.805	0.001
HRV_CSI_Modified	5.850	0.005
HRV_C2a	5.865	0.002
HRV_MFDFA_alpha1_Width	5.945	0.001
HRV_C1d	5.965	-0.002
HRV_Ca	5.965	0.001
HRV_MFDFA_alpha1_Delta	5.975	0.001
HRV_MFDFA_alpha1_Mean	6.000	-0.002
HRV_SD2	6.020	0.004
HRV_LFHF	6.025	0.001
HRV_MFDFA_alpha1_Increment	6.100	0.001
HRV_Cd	6.105	0.001
HRV_DFA_alpha1	6.125	-0.001
HRV_HFD	6.135	0.000

HRV_LFn	6.160	0.001
HRV_MSEn	6.190	0.000
HRV_CVSD	6.240	0.000
HRV_TP	6.295	0.001
HRV_SD1a	6.320	-0.001
HRV_SDNN	6.390	0.005
HRV_SD1SD2	6.395	0.002
HRV_HF	6.425	0.001
HRV_CVNN	6.430	0.001
HRV_LnHF	6.445	0.001
HRV_SDNNa	6.535	0.004
HRV_S	6.635	0.001
HRV_CSI	6.670	0.001
HRV_CVI	6.760	0.003
HRV_SDNNd	6.880	0.002
HRV_SD2d	6.920	0.001
HRV_SDSd	6.985	0.000
HRV_SD1d	7.010	0.001
HRV_SD1	7.020	0.000
HRV_RMSSD	7.195	0.000

```
stopCluster(cl)
```

```
vars_ranked <- rsf_var_select$topvars
```

Cross Validation to Select the Best Number of Features

We will use 10-fold cross validation to select the best number of features used in the model.

```
set.seed(1234)
folds <- createFolds(data$event, k = n_folds) # return indices of folds

cv_errors <- pmclapply(seq(1, length(vars_ranked), by = 1),
  ↪ function(num_vars) {
    selected_vars <- vars_ranked[1:num_vars]

    fold_errors <- sapply(folds, function(fold_idx) {
      train_data_fold <- data[-fold_idx, c("time", "event", selected_vars)]
      val_data_fold <- data[fold_idx, c("time", "event", selected_vars)]
```

```

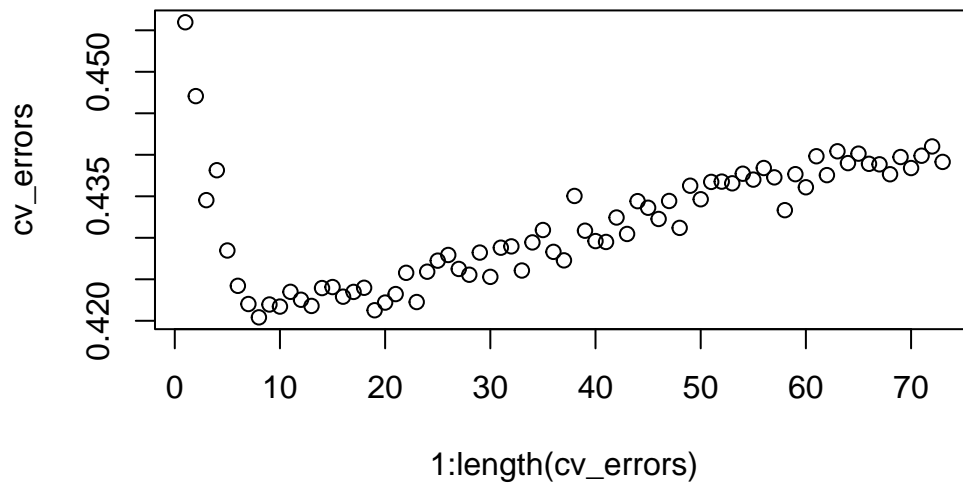
model <- rfsrc.fast(Surv(time, event) ~ .,
  data = train_data_fold,
  ntree = 200,
  forest = TRUE
)
pred <- predict(model,
  newdata = val_data_fold,
  na.action = "na.impute" # * There may be missing values in the
  ↪ dataset
)$predicted # define pred has attributes survival(sample_size*time)
↪ and predicted(sample_size) for risk
# Use C-index to measure the performance of the model
1 - concordance.index(
  pred, # pass risk prediction for first argument
  val_data_fold$time,
  val_data_fold$event
)$c.index
})
mean(fold_errors)
}, title = "Cross Validation to Select the Best Number of Features")

```

```

cv_errors <- as.numeric(cv_errors)
plot(1:length(cv_errors), cv_errors)

```

```
best_num_vars <- which.min(cv_errors)
vars_selected <- vars_ranked[1:best_num_vars]
```

```
print(paste0("The best number of features to retain is ", best_num_vars))
```

```
[1] "The best number of features to retain is 8"
```

Model Fitting

```
data_selected <- data[, c("time", "event", vars_selected)]

# Before formally fitting the model, we can tune the hyperparameters to find:
# 1. optimal mtry (possible split at each node)
# 2. optimal nodesize (minimum size of terminal nodes)
rsf_tuned <- tune.rfsrc(
  Surv(time, event) ~ .,
  data = data_selected,
)
```

```
rsf_model <- rfsrc(Surv(time, event) ~ .,  
  data = data_selected,  
  ntree = 500,  
  mtry = rsf_tuned$best.mtry,  
  nodesize = rsf_tuned$best.nodesize  
)
```

```
# We also fit the full model  
rsf_tuned_full <- tune.rfsrc(  
  Surv(time, event) ~ .,  
  data = data,  
)  
  
rsf_model_full <- rfsrc(Surv(time, event) ~ .,  
  data = data,  
  ntree = 1000,  
  mtry = rsf_tuned_full$best.mtry,  
  nodesize = rsf_tuned_full$best.nodesize  
)
```

```
# plot.rfsrc?  
# plot.variable.rfsrc?  
# print.rfsrc?
```